

Multivariate lookups based on logarithmic derivatives

Ulrich Haböck

Orbis Labs
team@orbislabs.com

December 12, 2022*

Abstract

Logarithmic derivatives translate products of linear factors into sums of their reciprocals, turning zeroes into simple poles of same multiplicity. Based on this simple fact, we construct an interactive oracle proof for lookups over the boolean hypercube, which makes use of a single multiplicity function instead of working with a rearranged union of table and witnesses, as Plookup [GW20] does. For single-column lookups the performance is comparable to Plookup, taken to the multivariate setting by Hyperplonk+ [CBBZ22]. However, the real power of our argument unfolds in the case of “batch-column” lookups, where multiple columns are subject to the same table lookup: While the number of field operations is comparable to the Plookup strategy, the oracles provided by our prover are significantly fewer. For example, given columns of length between 2^{12} and 2^{18} , and assuming a commitment scheme over an elliptic curve with a 256 bit large base field, our paper-pencil operation counts indicate that the logarithmic derivative lookup is more than 3.5 times faster.

Contents

1	Introduction	2
2	Preliminaries	3
2.1	The Lagrange kernel of the boolean hypercube	3
2.2	The formal derivate	4
2.3	The logarithmic derivative	5
2.4	Lagrange interactive oracle proofs	7
2.5	The sumcheck protocol	7
3	Lookups based on the logarithmic derivative	9
3.1	The protocol	10
3.2	Soundness	12
3.3	Computational cost and optimal batch size	12
3.4	The optimal batch size	13
3.5	Generalizations	14

*This updated version includes a discussion of the univariate variant of the argument.

4	Lookups based on the Hyperplonk shift	15
4.1	Multivariate Plookup	15
4.2	A reference variant for Section 4.1	16
4.3	Computational cost	17
4.4	Comparison with logarithmic derivative lookups	18
5	Acknowledgements	18
A	Appendix: the univariate case	21
A.1	The logarithmic derivate approach	22
A.2	Plookup	24
A.3	The flookup strategy	25
A.4	Comparison	27

1 Introduction

Lookup arguments prove a sequence of values being member of an, often prediscrbed, table. They are an essential tool for improving the efficiency of arguments for statements which are otherwise quite expensive to arithmetize. Main applications are (1) lookups for relations of high algebraic complexity, and (2) lookups for interval ranges, the latter of which are extensively used by zero-knowledge virtual machines. Although closely related to permutation (or shuffle) arguments [BG12], a first explicit occurence of lookups dates back to Arya [BCG⁺18], to the best of our knowledge. While their argument handles multiplicities directly in a quite costly manner, Plookup [GW20] greatly improved over [BCG⁺18] using a rather geometric approach. Since then Plookup (and variants of it) is *the* general purpose lookup argument used in practical applications.

In this paper we describe a lookup argument based which is based on logarithmic derivatives. As in classical calculus, formal logarithmic derivatives turn products $\prod_{i=1}^N (X - z_i)$ into sums of their reciprocals,

$$\sum_{i=1}^N \frac{1}{X - z_i},$$

having poles with the same multiplicity as the zeros of the product. Working with poles instead of zeros is extremly useful for lookup arguments. While the treatment of multiplicities is costly in the product approach, it turns cheap when using logarithmic derivatives: Given a sequence of field elements $(a_i)_{i=1}^N$ and another sequence $(t_j)_{j=1}^M$, then $\{a_i : i = 1, \dots, N\} \subseteq \{t_j : j = 1, \dots, M\}$ as sets, if and only if there exists a sequence of field elements $(m_j)_{j=1}^M$ (the multiplicities) such that

$$\sum_{i=1}^N \frac{1}{X - a_i} = \sum_{j=1}^M \frac{m_j}{X - t_j}.$$

(This holds under quite mild conditions on the field, see Lemma 5 for details.) Based on this fractional identity we construct a lookup proof which is more efficient than the Plookup approach, which argues via a sorted union of witness and table sequence. This is particularly true in the case of *batch-column* lookups, where several sequences (“columns”) are subject to the same table lookup, a situation that often arises in zero-knowledge virtual machines, enforcing execution trace elements being valid machine words. For example, the arithmetic unit of Polygon’s System Zero [Pol] has about 70 columns subject to one and the same range check, or the tinyRAM implementation of Orbis Labs [Lab] has 12 columns subject to one and

the same lookup table. In our lookup the oracle costs, measuring the number and sizes of the oracles, are significantly lower than for a lookup based on the Plookup strategy.

We stress the fact that we are not the only ones who exploit fractional decompositions for lookups. Concurrently, building on the work of [ZBK⁺22] and [PK22] on “large-table” lookups, Gabizon and Khovratovich [GK22] describe a bilinear argument, the proving-time of which is independent of the table size. The main contribution of [GK22] is a univariate oracle proof for the radical of a witness sequence (i.e. the sequence with multiple occurrences removed) which is almost identical to our approach¹. We give comparison with our approach in the appendix.

This document focuses on batch-column lookups with asymptotically linear prover costs. It is organized as follows. In Section 2, we gather the preliminaries used in the sequel: The Lagrange kernel over the boolean hypercube, basic facts on the formal logarithmic derivative, and a summary of the multivariate sumcheck argument. Besides that, we informally introduce Lagrange interactive oracle proofs, an oracle model we consider suitable for arguments which are based on the Lagrange representation of polynomials rather than their coefficients. In Section 3 we describe our lookup argument based on the logarithmic derivative. For comparison reasons, we add an extra section (Section 4) in which we outline batch-column lookups using the Plookup strategy, adapted to the boolean hypercube. These rely on the time shift from Hyperplonk [CBBZ22], and we consider them state-of-the-art in the multivariate setting.

We finally point out, that although our protocol is written for the multilinear setting, its translation into a univariate proof is straight-forward. However, in the univariate context the benefit of logarithmic derivatives is much less. This is essentially due to the demand of overall quotients with respect to the witness domain, the degree of which increases with the batch size. In this updated version of the document, we give a detailed treatise of the univariate protocol in Appendix A.

2 Preliminaries

2.1 The Lagrange kernel of the boolean hypercube

Let F denote a finite field, and F^* its multiplicative group. Throughout the document we regard the boolean hypercube $H = \{\pm 1\}^n$ as a multiplicative subgroup of $(F^*)^n$. For a multivariate function $f(X_1, \dots, X_n)$, we will often use the vector notation $\vec{X} = (X_1, \dots, X_n)$ for its arguments, writing $f(\vec{X}) := f(X_1, \dots, X_n)$.

The *Lagrange kernel* of H is the multilinear polynomial

$$L_H(\vec{X}, \vec{Y}) = \frac{1}{2^n} \cdot \prod_{j=1}^n (1 + X_j \cdot Y_j). \quad (1)$$

Notice that $L_H(\vec{X}, \vec{Y})$ is symmetric in \vec{X} and \vec{Y} , i.e. $L_H(\vec{X}, \vec{Y}) = L_H(\vec{Y}, \vec{X})$, and that (1) is evaluated within only $\mathcal{O}(\log |H|)$ field operations. Whenever $\vec{y} \in H$ we have that $L_H(\vec{X}, \vec{y})$ is the Lagrange polynomial on H , which is the unique multilinear polynomial which satisfies $L_H(\vec{x}, \vec{y}) = 1$ at $\vec{x} = \vec{y}$, and zero elsewhere on H . In particular for a function $f : H \rightarrow F$ the inner product evaluation formula

$$\langle f, L_H(\cdot, \vec{y}) \rangle_H := \sum_{\vec{x} \in H} f(\vec{x}) \cdot L_H(\vec{x}, \vec{y}) = f(\vec{y}).$$

is valid for every $\vec{y} \in H$. This property extends beyond H , as the following Lemma shows.

¹Almost simultaneously, G. Roh, W. Dai, M. Jabbour, and A. He published the same idea in a blog post [RDJH22]. To underpin that our work was not influenced by concurrent efforts we refer to the verified commit on our Github repository, <https://github.com/Orbis-Tertius/MVlookups/commit/4a8816a04e8107e05d4bb0897ee52e72210c84b8>

Lemma 1. Let $p(\vec{X})$ be the unique multilinear extension of $f : H \rightarrow F$. Then for every $\vec{y} \in F^n$,

$$\langle f, L_H(\cdot, \vec{y}) \rangle_H = \sum_{x \in H} f(\vec{x}) \cdot L_H(\vec{x}, \vec{y}) = p(\vec{y}). \quad (2)$$

Proof. Since $p(\vec{y}) = \sum_{\vec{x} \in H} f(\vec{x}) \cdot L_H(\vec{X}, \vec{x})$, it suffices to show the claim for $p(X) = L_H(\vec{X}, \vec{z})$, with $\vec{z} \in H$. By the property of $L_H(\vec{X}, \vec{z})$, we have $\langle L_H(\cdot, \vec{z}), L_H(\cdot, \vec{y}) \rangle_H = L_H(\vec{y}, \vec{z})$, which by symmetry is equal to $L_H(\vec{X}, \vec{y})$ at $\vec{X} = \vec{z}$. This completes the proof of the Lemma. \square

Note that for any $\vec{y} \in F^n$, the domain evaluation of $L_H(\vec{X}, \vec{y})$ over H can be computed in $\mathcal{O}(|H|)$ field operations, by recursively computing the domain evaluation of the partial products $p_k(X_1, \dots, X_k, y_1, \dots, y_k) = \frac{1}{2^n} \cdot \prod_{j=1}^k (1 + X_j \cdot y_j)$ over $H_k = \{\pm 1\}^k$ from the domain evaluation of p_{k-1} , where one starts with $f_0 = \frac{1}{2^n}$ over the single-point domain H_0 . Each recursion step costs $|H_{k-1}|$ field multiplications, denoted by M , and the same number of additions, denoted by A , yielding overall

$$\sum_{k=1}^n |H_{k-1}| \cdot (M + A) < |H| \cdot (M + A). \quad (3)$$

2.2 The formal derivate

Given a univariate polynomial $p(X) = \sum_{k=0}^d c_k \cdot X^k$ over a general (possibly infinite) field F , its *derivative* is defined as

$$p'(X) := \sum_{k=1}^d k \cdot c_k \cdot X^{k-1}. \quad (4)$$

As in calculus, the derivative is linear, i.e. for every two polynomials $p_1(X), p_2(X) \in F[X]$, and coefficients $\lambda_1, \lambda_2 \in F$,

$$(\lambda_1 \cdot p_1(X) + \lambda_2 \cdot p_2(X))' = \lambda_1 \cdot p_1'(X) + \lambda_2 \cdot p_2'(X)$$

and we have the product rule

$$(p_1(X) \cdot p_2(X))' = p_1'(X) \cdot p_2(X) + p_1(X) \cdot p_2'(X).$$

For a function $\frac{p(X)}{q(X)}$ from the rational function field $F(X)$, the derivative is defined as the rational function

$$\left(\frac{p(X)}{q(X)} \right)' := \frac{p'(X) \cdot q(X) - p(X) \cdot q'(X)}{q(X)^2}. \quad (5)$$

By the product rule for polynomials, the definition does not depend on the representation of $\frac{p(X)}{q(X)}$. Both linearity as well as the product rule extend to rational functions.

For any polynomial $p(X) \in F[X]$, if $p'(X) = 0$ then $p(X) = g(X^p)$ where p is the characteristic of the field F . In particular, if $\deg p(X) < p$, then the polynomial must be constant. As the analogous fact for fractions is not as commonly known, we give a proof of the following lemma.

Lemma 2. Let F be a field of characteristic $p \neq 0$, and $\frac{p(X)}{q(X)}$ a rational function over F with both $\deg p(X) < p$ and $\deg q(X) < p$. If the formal derivative $\left(\frac{p(X)}{q(X)} \right)' = 0$, then $\frac{p(X)}{q(X)} = c$ for some constant $c \in F$.

Proof. If $q(X)$ is a constant, then the assertion of the Lemma follows from the corresponding statement for polynomials. Hence we assume that $\deg q(X) > 0$. Use polynomial division to obtain the representation

$$\frac{p(X)}{q(X)} = m(X) + \frac{r(X)}{q(X)},$$

with $m(X), r(X) \in F[X]$, $\deg m(X) \leq \deg p(X)$, and $\deg r(X) < \deg q(X)$ whenever $r(X) \neq 0$. By linearity of the derivative, we have $0 = \left(\frac{p(X)}{q(X)}\right)' = m'(X) + \left(\frac{r(X)}{q(X)}\right)'$, and therefore

$$r'(X) \cdot q(X) - r(X) \cdot q'(X) = -m'(X) \cdot q(X)^2. \quad (6)$$

Comparing the degrees of left and right hand side in (6), we conclude that $m'(X) = 0$. Since $\deg m(X) \leq \deg p(X) < p$ we have $m(X) = c$ for some constant² $c \in F$. Furthermore, if we had $r(X) \neq 0$ then the leading term of the left hand side in (6) would be

$$(k - n) \cdot c_n \cdot d_k \cdot X^{n+k-1},$$

with $c_n \cdot X^n$, $n > 0$, being the leading term of $q(X)$, and $d_k \cdot X^k$, $0 \leq k < n$, the leading term of $r(X)$. As $0 < n - k < p$, and both $c_n \neq 0$ and $c_m \neq 0$, the leading term of the left hand side of (6) would not vanish. Therefore it must hold that $r(X) = 0$ and the proof of the lemma is complete. \square

2.3 The logarithmic derivative

The *logarithmic derivate* of a polynomial $p(X)$ over a (general) field F is the rational function

$$\frac{p'(X)}{p(X)}.$$

Note that the logarithmic derivative of the product $p_1(X) \cdot p_2(X)$ of two polynomials $p_1(X), p_2(X)$ equals the sum of their logarithmic derivatives, since by the product rule we have

$$\frac{(p_1(X) \cdot p_2(X))'}{p_1(X) \cdot p_2(X)} = \frac{p_1'(X) \cdot p_2(X) + p_1(X) \cdot p_2'(X)}{p_1(X) \cdot p_2(X)} = \frac{p_1'(X)}{p_1(X)} + \frac{p_2'(X)}{p_2(X)}.$$

In particular the logarithmic derivative of a product $p(X) = \prod_{i=1}^n (X + z_i)$, with each $z_i \in F$, is equal to the sum

$$\frac{p'(X)}{p(X)} = \sum_{i=1}^n \frac{1}{X + z_i}. \quad (7)$$

The following lemma³ is a simple consequence of Lemma 2 and essentially states that, under quite mild conditions on the field F , if two normalized polynomials have the same logarithmic derivative then they are equal. We state this fact for our use case of product representations.

Lemma 3. *Let $(a_i)_{i=1}^n$ and $(b_i)_{i=1}^n$ be sequences over a field F with characteristic $p > n$. Then $\prod_{i=1}^n (X + a_i) = \prod_{i=1}^n (X + b_i)$ in $F[X]$ if and only if*

$$\sum_{i=1}^n \frac{1}{X + a_i} = \sum_{i=1}^n \frac{1}{X + b_i}$$

in the rational function field $F(X)$.

²For general degrees of $p(X)$ we would only be able to conclude that $m(X) = g(X^p)$ for some polynomial $g(X)$.

³At the time of writing we learned that also others are aware of Lemma 3. In a Delendum Frontiers session on multi-party computation (6. Oct. 2022) an audience member pointed out its usefulness for permutation arguments.

Proof. If $p_a(X) = \prod_{i=1}^n (X + a_i)$ and $p_b(X) = \prod_{i=1}^n (X + b_i)$ coincide, so do their logarithmic derivatives. To show the other direction, assume that $\frac{p'_a(X)}{p_a(X)} = \frac{p'_b(X)}{p_b(X)}$. Then

$$\left(\frac{p_a(X)}{p_b(X)}\right)' = \frac{p'_a(X) \cdot p_b(X) - p_a(X) \cdot p'_b(X)}{p_b^2(X)} = 0.$$

Hence by Lemma 2 we have $\frac{p_a(X)}{p_b(X)} = c$ for some constant $c \in F$. As both $p_a(X)$ and $p_b(X)$ have leading coefficient equal to 1, we conclude that $c = 1$, and the proof of the Lemma is complete. \square

Remark 1. We stress the fact that Lemma 3 also applies to the case where F is the function field $F_p(Y_1, \dots, Y_k)$ over a finite field F_p of characteristic p . This observation will be useful when generalizing the permutation argument to the case where a_i and b_i are multilinear polynomials in Y_1, \dots, Y_n .

Given a product $p(X) = \prod_{i=1}^N (X + a_i)$ we can gather the poles of its logarithmic derivative obtaining the fractional decomposition

$$\frac{p'(X)}{p(X)} = \sum_{a \in F} \frac{m(a)}{X + a},$$

where $m(a) \in \{1, \dots, N\}$ is the multiplicity of the value a in $(a_i)_{i=1}^N$. Fractional decompositions are unique, as shown by the following lemma.

Lemma 4. *Let F be an arbitrary field and $m_1, m_2 : F \rightarrow F$ any functions. Then $\sum_{z \in F} \frac{m_1(z)}{X - z} = \sum_{z \in F} \frac{m_2(z)}{X - z}$ in the rational function field $F(X)$, if and only if $m_1(z) = m_2(z)$ for every $z \in F$.*

Proof. Suppose that the fractional decompositions are equal. Then $\sum_{z \in F} \frac{m_1(z) - m_2(z)}{X - z} = 0$, and therefore

$$p(X) = \prod_{w \in F} (X - w) \cdot \sum_{z \in F} \frac{m_1(z) - m_2(z)}{X - z} = \sum_{z \in F} (m_1(z) - m_2(z)) \cdot \prod_{w \in F \setminus \{z\}} (X - w) = 0.$$

In particular, $p(z) = (m_1(z) - m_2(z)) \cdot \prod_{w \in F \setminus \{z\}} (z - w) = 0$ for every $z \in F$. Since $\prod_{w \in F \setminus \{z\}} (z - w) \neq 0$ we must have $m_1(z) = m_2(z)$ for every $z \in F$. The other direction is obvious. \square

This leads to the following algebraic criterion for set membership, which is the key tool for our lookup arguments.

Lemma 5 (Set inclusion). *Let F be a field of characteristic $p > N$, and suppose that $(a_i)_{i=1}^N, (b_i)_{i=1}^N$ are arbitrary sequences of field elements. Then $\{a_i\} \subseteq \{b_i\}$ as sets (with multiples of values removed), if and only if there exists a sequence $(m_i)_{i=1}^N$ of field elements from $F_q \subseteq F$ such that*

$$\sum_{i=1}^N \frac{1}{X + a_i} = \sum_{i=1}^N \frac{m_i}{X + b_i} \tag{8}$$

in the function field $F(X)$. Moreover, we have equality of the sets $\{a_i\} = \{b_i\}$, if and only if $m_i \neq 0$, for every $i = 1, \dots, N$.

Proof. Let us denote by $m_a(z)$ the multiplicity of a field element z in the sequence $(a_i)_{i=1}^N$. Likewise, we do for $(b_i)_{i=1}^N$. Note that since $N < p$, the multiplicities can be regarded as non-zero elements from F_p as a subset of F . Suppose that $\{a_i\} \subseteq \{b_i\}$ as sets. Set (m_i) as the normalized multiplicities $m_i = \frac{m_a(b_i)}{m_b(b_i)}$. This choice of (m_i) obviously satisfies (8).

Conversely, suppose that (8) holds. Collecting fractions with the same denominator we obtain fractional representations for both sides of the equation (8),

$$\begin{aligned}\sum_{i=1}^N \frac{1}{X + a_i} &= \sum_{z \in F} \frac{m_a(z)}{X + z}, \\ \sum_{i=1}^N \frac{m_i}{X + b_i} &= \sum_{z \in F} \frac{\mu(z)}{X + z}.\end{aligned}$$

Note that since $N < p$, we know that for each $z \in \{a_i\}$ we have $m_a(z) \neq 0$. By the uniqueness of fractional representations, Lemma 4, $m_a(z) = \mu(z)$ for every $z \in \{a_i\}$, and therefore each $z \in \{a_i\}$ must occur also in $\{b_i\}$. \square

2.4 Lagrange interactive oracle proofs

The oracle proofs of many general purpose SNARKs such as Plonk [GWC19] or algebraic intermediate representations [BSBHR18] rely on witnesses that are given in Lagrange representation, i.e. by their values over a domain H . Their multivariate variants may completely avoid the usage of fast Fourier transforms whenever the polynomial commitment scheme can be turned into one that does not need to know the coefficients, neither when computing a commitment nor in an opening proof. Exactly this property is captured by *Lagrange oracle proofs*, rather than polynomial ones [BFS20].

A *Lagrange interactive oracle proof* (*Lagrange IOP*) over the boolean hypercube $H = \{\pm 1\}^n$ is an interactive protocol between two parties, the “prover” and the “verifier”. In each round, the verifier sends a message (typically a random challenge) and the prover computes one or several functions over the boolean hypercube, and gives the verifier oracle access to them. From the moment on it is given access, the verifier is allowed to query the oracles for their inner products with the Lagrange kernel $L_H(\cdot, \vec{y})$, associated with an arbitrary vector $\vec{y} \in F^n$.

The security notions for Lagrange IOPs, such as completeness, (knowledge) soundness, and zero-knowledge, are exactly the same as for other interactive oracle proofs. We assume that the reader is familiar with these, and refer to [BSCS16] or [BFS20] for their formal definitions.

Lagrange IOPs are turned into arguments by instantiating the Lagrange oracles by a *Lagrange commitment scheme*. A Lagrange commitment scheme is a commitment scheme for functions over H that comes with an evaluation proof for Lagrange queries. For example, inner product arguments [BCC⁺16] can be directly used to construct Lagrange commitment schemes, but also the multilinear variant [PST13] of the [KZG10] commitment scheme is easily modified to completely avoid dealing with coefficients. We suppose that this is well-known, and therefore we omit an explicit elaboration in this paper.

2.5 The sumcheck protocol

We give a concise summary on the multivariate sumcheck protocol [LFKN92]. Given a multivariate polynomial $p(X_1, \dots, X_n) \in F[X_1, \dots, X_n]$, a prover wants to convince a verifier upon that

$$s = \sum_{(x_1, \dots, x_n) \in \{\pm 1\}^n} p(x_1, \dots, x_n).$$

This is done by a random folding procedure which, starting with $H_0 = \{\pm 1\}^n$, which stepwise reduces a claim on the sum over $H_i = \{\pm 1\}^{n-i}$, $i = 0, \dots, n-1$, to one over the hypercube H_{i+1} of half the size. Eventually, one ends up with a claim over a single-point sum, which is paraphrased as the value of $p(X_1, \dots, X_n)$ at a random point $(r_1, \dots, r_n) \in F^n$ sampled in the course of the reduction steps.

Protocol 1 (Sumcheck protocol, [LFKN92]). Let $p(X_1, \dots, X_n)$ be a multivariate polynomial over a finite field F . The sumcheck protocol, in which a prover wants to convince the verifier upon the sum $s = \sum_{(x_1, \dots, x_n) \in \{\pm 1\}^n} p(x_1, \dots, x_n)$, is as follows. We write $s_0(X)$ for the constant polynomial $s_0 = s$.

- In each round $i = 1, \dots, n$, the prover sends the coefficients of the univariate polynomial

$$s_i(X) = \sum_{(x_{i+1}, \dots, x_n) \in \{\pm 1\}^{n-i}} p(r_1, \dots, r_{i-1}, X, x_{i+1}, \dots, x_n),$$

of degree $d_i \leq \deg_{X_i} p(X_1, \dots, X_n)$, where r_1, \dots, r_{i-1} are the randomnesses received in the previous rounds. (In the first round $i = 1$ there are no previous randomnesses, and $p(r_1, \dots, r_{i-1}, X, x_{i+1}, \dots, x_n)$ is meant to denote $p(X, x_2, \dots, x_n)$.) The prover sends the coefficients of $s_i(X)$ to the verifier, which checks whether the received polynomial $s_i(X)$ is in fact of the expected degree and that

$$s_{i-1}(r_{i-1}) = s_i(+1) + s_i(-1).$$

(Again, in the first round $i = 1$ there is no r_0 , and the verifier checks wheather $s_0 = s_1(+1) + s_1(-1)$.) If so, the verifier samples random challenge $r_i \leftarrow_{\$} F$ uniformly from F and sends it to the prover.

After these rounds the verifier checks that $s_n(r_n) = p(r_1, \dots, r_n)$. If so, the verifier accepts (otherwise it rejects).

Soundness of the sumcheck protocol is proven by a repeated application of the Schwartz-Zippel lemma. We omit a proof, and refer to [LFKN92] or [Tha13].

Theorem 2 ([LFKN92]). The sumcheck protocol (Protocol 1) has soundness error

$$\varepsilon_{\text{sumcheck}} \leq \frac{1}{|F|} \cdot \sum_{i=1}^n \deg_{X_i} p(X_1, \dots, X_n). \quad (9)$$

The sumcheck protocol is easily extended to the sumcheck for a batch of polynomials $p_i(X_1, \dots, X_n)$, $i = 0, \dots, L$, by letting the verifier sample a random vector $(\lambda_1, \dots, \lambda_L) \leftarrow_{\$} F^L$, and a subsequent sumcheck protocol for the random linear combination

$$\bar{p}(X_1, \dots, X_n) = p_0(X_1, \dots, X_n) + \sum_{i=1}^L \lambda_i \cdot p_i(X_1, \dots, X_n).$$

The soundness error bound increases only slightly,

$$\varepsilon_{\text{sumcheck}} \leq \frac{1}{|F|} \cdot \left(1 + \sum_{i=1}^n \deg_{X_i} p(X_1, \dots, X_n) \right). \quad (10)$$

Computational cost

Let us discuss the prover cost of the sumcheck protocol for the case that $p(\vec{X}) = p(X_1, \dots, X_n)$ is of the form

$$p(\vec{X}) = Q(w_1(\vec{X}), \dots, w_m(\vec{X})),$$

with each $w_i(\vec{X}) \in F[X_1, \dots, X_n]$ being multilinear, and

$$Q(Y_1, \dots, Y_m) = \sum_{(i_1, \dots, i_m) \in \{0,1\}^m} c_{i_1, \dots, i_m} \cdot Y_1^{i_1} \cdots Y_m^{i_m}$$

is a multivariate polynomial having (a typically low) absolute degree d . We denote the arithmetic complexity, i.e. the number of field multiplications M , subtractions S and additions A to evaluate Q by $|Q|_M$, $|Q|_S$ and $|Q|_A$, respectively. Each of the univariate polynomials $s_i(X)$, $i = 1, \dots, n$, is of degree at most d the absolute degree of Q , and is computed from its values over a set $D \supseteq \{\pm 1\}$ of size $|D| = d + 1$. In each step $i = 1, \dots, n$, the values of $s_i(z)$ for $z \in D$ are obtained by linear interpolation of the domain evaluations of each

$$w_j(r_1, \dots, r_{i-1}, \pm 1, X_{i+1}, \dots, X_n)$$

over $H_i = \{\pm 1\}^{n-i}$ as given from the previous step, to the domain evaluation

$$w_j(r_1, \dots, r_{i-1}, z, X_{i+1}, \dots, X_n),$$

the values of which are used for computing $s_i(z) = \sum_{(x_{i+1}, \dots, x_n) \in H_i} Q(r_1, \dots, r_{i-1}, z, x_{i+1}, \dots, x_n)$. Given the random challenge r_i from the verifier, the domain evaluation of each

$$w_j(r_1, \dots, r_{i-1}, r_i, X_{i+1}, \dots, X_n)$$

is computed by another linear interpolation. Linear interpolation costs $|H_i|$ multiplications and the same number of additions/subtractions for each multilinear polynomial, the values of Q are obtained within $|Q|_M \cdot M + |Q|_S \cdot S + |Q|_A \cdot A$. In terms of field multiplications M , subtractions S and additions A , step i consumes

$$m \cdot |H_i| \cdot S + m \cdot (|D| - 1) \cdot |H_i| \cdot (M + A) + |D| \cdot |H_i| \cdot (|Q|_M \cdot M + |Q|_S \cdot S + |Q|_A \cdot A) + |D| \cdot |H_i| \cdot A,$$

where the last term is for the domain sums. Since $\sum_{i=1}^n |H_i| = |H| - 1$, the overall cost for the prover is bounded by

$$|H| \cdot \left(1 - \frac{1}{|H|}\right) \cdot ((d \cdot m + (d+1) \cdot |Q|_M) \cdot M + (m + (d+1) \cdot |Q|_S) \cdot S + (d \cdot m + (d+1) \cdot (|Q|_A + 1)) \cdot A). \quad (11)$$

We shall use this formula for the operation counts of our lookup protocol.

3 Lookups based on the logarithmic derivative

Assume that F is a finite field, and that f_1, \dots, f_M and $t : H \rightarrow F$ are functions over the Boolean hypercube $H = \{\pm 1\}^n$. By Lemma 5, it holds that $\bigcup_{i=1}^M \{f_i(\vec{x})\}_{\vec{x} \in H} \subseteq \{t(\vec{x})\}_{\vec{x} \in H}$ as sets, if and only if there exists a function $m : H \rightarrow F$ such that

$$\sum_{\vec{x} \in H} \sum_{i=1}^M \frac{1}{X + f_i(\vec{x})} = \sum_{\vec{x} \in H} \frac{m(\vec{x})}{X + t(\vec{x})}, \quad (12)$$

assuming that the characteristic of F is larger than M times the size of the hypercube. If t is injective (which is typically the case for lookup tables) then m is the multiplicity function, counting the number of occurrences for each value $t(\vec{x})$ in f_1, \dots, f_M altogether, i.e. $m(\vec{x}) = m_f(t(\vec{x})) = \sum_{i=1}^M |\{\vec{y} \in H : f_i(\vec{y}) = t(\vec{x})\}|$. If t is not one-to-one, we set m as the *normalized* multiplicity function

$$m(\vec{x}) = \frac{m_f(t(\vec{x}))}{m_t(t(\vec{x}))} = \frac{\sum_{i=1}^M |\{\vec{y} \in H : f_i(\vec{y}) = t(\vec{x})\}|}{|\{\vec{y} \in H : t(\vec{y}) = t(\vec{x})\}|}. \quad (13)$$

Given a random challenge $x \leftarrow F$ from the verifier, the prover shows that the rational identity (12) holds at $X = x$,

$$\sum_{\vec{x} \in H} \sum_{i=1}^M \frac{1}{x + f_i(\vec{x})} - \frac{m(\vec{x})}{x + t(\vec{x})} = 0,$$

whenever evaluation is possible. However, in order to apply the sumcheck protocol we need to turn the fractional expression into a polynomial one. For that, one may split the sum into partial ones (of equal size) and provide multilinear helper functions for each sum. These helper functions are subject to a domain identity over H , the algebraic degree of which corresponds to the number of reciprocals in the sum. The optimal size of the partial sums depends on the used polynomial commitment scheme. The costlier a commitment, the higher the algebraic degree for the domain identity of the helper function can be.

Our building block lookup protocol, Protocol 2, saves as much commitments as possible. The prover provides a single multilinear helper function for the entire sumcheck expression, at the cost of an overall domain identity of absolute degree M , roughly. In Section 3.2 we sketch its soundness, and in Section 3.4 we give detailed operation counts for the oracle prover. We then estimate M_{opt} the optimal number of columns for Protocol 2, assuming a commitment scheme in an ordinary sized elliptic curve⁴. We stress the fact that our notion of optimality refers to the minimal cost per column of Protocol 2, which is a somewhat simplified cost measure. It is accurate when processing larger batches of columns by parallel invocations of the protocol on M_{opt} columns each, a strategy which we consider for simplicity of the presentation. In practice one uses a single multiplicity function and a single sumcheck across several invocations of the protocol.

3.1 The protocol

The prover provides the normalized multiplicity function m as defined in (13), and the helper function

$$h(\vec{x}) = \sum_{i=1}^M \frac{1}{x + f_i(\vec{x})} - \frac{m(\vec{x})}{x + t(\vec{x})}, \quad (14)$$

subject to the sumcheck $\sum_{\vec{x} \in H} h(\vec{x}) = 0$. Correctness of h is ensured by the domain identity

$$(h(\vec{x}) \cdot (x + t(\vec{x})) + m(\vec{x})) \cdot \prod_{i=1}^M (x + f_i(\vec{x})) = (x + t(\vec{x})) \cdot \sum_{i=1}^M \prod_{j \neq i} (x + f_j(\vec{x})) \quad (15)$$

over H , which is reduced to another sumcheck by applying the Lagrange kernel $L_H(\cdot, \vec{z})$ at a randomly chosen $\vec{z} \leftarrow_{\$} F^n$. Both sumchecks, the one for h and the one for the domain identity, are then combined into a single one, using another randomness $\lambda \leftarrow_{\$} F$.

Protocol 2 (Batch-column lookup over $H = \{\pm 1\}^n$). *Let $M \geq 1$ be an integer, and F a finite field with characteristic $p > M \cdot 2^n$. Given any functions $f_1, \dots, f_M, t : H \rightarrow F$ on the boolean hypercube $H = \{\pm 1\}^n$, the Lagrange IOP for that $\bigcup_{i=1}^M \{f_i(\vec{x}) : \vec{x} \in H\} \subseteq \{t(\vec{x}) : \vec{x} \in H\}$ as sets is as follows.*

1. *The prover determines the (normalized) multiplicity function $m : H \rightarrow F$ as defined in (13), and sends the oracle for m to the verifier. The verifier answers with a random sample $x \leftarrow_{\$} F \setminus \{-t(\vec{x}) : \vec{x} \in H\}$.*
2. *Given the challenge x from the verifier, the prover computes the randomized functions $\varphi_i(\vec{x}) = x + f_i(\vec{x})$, $i = 1, \dots, M$, and $\tau(\vec{x}) = x + t(\vec{x})$. It determines the values for*

$$h(\vec{x}) = \sum_{i=1}^M \frac{1}{\varphi_i(\vec{x})} - \frac{m(\vec{x})}{\tau(\vec{x})}, \quad (16)$$

over H , and sends the oracle for h to the verifier.

⁴By ordinary sized we mean a 128 bit secure curve over a 256 bit large base field. This covers both KZG-like commitments in Barreto-Naehrig curves, and IPA commitments in an ordinary elliptic curve.

3. The verifier responds with a random vector $\vec{z} \leftarrow_{\$} F^n$ and a batching randomness $\lambda \leftarrow_{\$} F$. Now, both prover and verifier engage in the sumcheck protocol (Protocol 1) for

$$\sum_{\vec{x} \in H} Q(L_H(\vec{x}, \vec{z}), h(\vec{x}), m(\vec{x}), \varphi_1(\vec{x}), \dots, \varphi_M(\vec{x}), \tau(\vec{x})) = 0,$$

where

$$Q(L, h, m, \varphi_1, \dots, \varphi_M, \tau) = L \cdot \left((h \cdot \tau + m) \cdot \prod_{i=1}^M \varphi_i - \tau \cdot \sum_{i=1}^M \prod_{j \neq i} \varphi_j \right) + \lambda \cdot h. \quad (17)$$

The sumcheck protocol outputs the expected value v for the multivariate polynomial

$$Q(L_H(\vec{X}, \vec{z}), h(\vec{X}), m(\vec{X}), \varphi_1(\vec{X}), \dots, \varphi_M(\vec{X}), \tau(\vec{X})) \quad (18)$$

at $\vec{X} = \vec{r}$ sampled by the verifier in the course of the protocol.

4. The verifier queries $[f_1], \dots, [f_M], [t], [m], [h]$ for their inner product with $L_H(\cdot, \vec{r})$, and uses the answers to check whether (18) equals the expected value v at $\vec{X} = \vec{r}$. (The value $L_H(\vec{r}, \vec{z})$ is computed by the verifier.)

Remark 3. We imposed the condition $x \notin \{-t(\vec{x})\}_{\vec{x} \in H}$ merely for completeness. However in some applications it may be not be desirable, or even not possible, to sample x from outside the range of t . There are several ways to handle this. One can simply omit the constraint on x , letting the verifier sample $x \leftarrow_{\$} F$ and the prover set h arbitrary whenever (16) is not defined. This comes at no extra cost, but the obtained protocol is only overwhelmingly complete. That is, with a probability of at most $\frac{|H|}{|F|}$ in the verifier randomness x , the honest prover does not succeed. In practice this is often considered acceptable, and many lookup implementations have a non-zero completeness error. Whenever this is not acceptable, one may modify the domain identity (15) to

$$\left((h \cdot \tau + m) \cdot \prod_{i=1}^M \varphi_i - \tau \cdot \sum_{i=1}^M \prod_{j \neq i} \varphi_j \right) \cdot \tau \cdot \prod_{i=1}^M \varphi_i = 0 \quad (19)$$

over H , which imposes no condition on $h(\vec{x})$ whenever $\tau(\vec{x}) = 0$. However, this approach comes at the cost of almost doubling the absolute degree of Q .

Let us point out two variations of Protocol 2. In the single-column case $M = 1$ the lookup argument can be turned into a multiset check for the ranges of f_1 and t , by setting m as the constant function $m(\vec{x}) = 1$. In this case only h needs to be provided by the prover. More interestingly, Protocol 2 is easily extended to a proof of range equality, showing that $\bigcup_{i=1}^M \{f_i(\vec{x})\}_{\vec{x} \in H} = \{\tau(\vec{x})\}_{\vec{x} \in H}$ as sets. For this the prover additionally shows that $m \neq 0$ over H , which is done by providing another auxiliary function $h_m : H \rightarrow F$ subject to $h_m \cdot m = 1$ over H . However, we are not aware of any application of this fact.

3.2 Soundness

The soundness analysis of Protocol 2 is a straight-forward application of the Schwartz-Zippel lemma and the Lagrange-query to point-query correspondence stated by Lemma 1. We merely sketch it. The univariate rational lookup identity (12) is turned into a polynomial identity of degree at most $|H| \cdot (M + 1) - 1$ by multiplying it with the common denominator

$$p(X) = \prod_{\vec{x} \in H} (X + t(\vec{x})) \cdot \prod_{i=1}^M (X + f_i(\vec{x})). \quad (20)$$

Since we sample x from a set of size at least $|F| - |H|$, the soundness error of Step 2 of the protocol is at most

$$\varepsilon_1 \leq \frac{(M+1) \cdot |H| - 1}{|F| - |H|}. \quad (21)$$

The soundness error due to the reduction of the domain identity (15) to the Lagrange kernel based sumcheck is

$$\varepsilon_2 \leq \frac{1}{|F|},$$

as scalar products with the Lagrange kernel translate to point evaluation of the multilinear extension. This yields the following theorem.

Theorem 4. *The interactive oracle proof described Protocol 2 has soundness error*

$$\varepsilon < \frac{(M+1) \cdot |H| - 1}{|F| - |H|} + \varepsilon_{\text{sumcheck}},$$

where $\varepsilon_{\text{sumcheck}}$ is the soundness error of the sumcheck argument (10) over H for a multivariate polynomial in $M+4$ variables with maximum individual degree $M+3$.

3.3 Computational cost and optimal batch size

Let us determine the prover cost of Protocol 2. The polynomial Q from (17) has $\nu = M+4$ variables, and absolute degree $d = M+3$. Let us describe a domain evaluation strategy for the values of Q , which makes use of batch inversion. This strategy allows us to evaluate Q much more efficiently than using (17), but demands a modification of the sumcheck operation count formula (11). Assume that the inverses of $\varphi_1, \dots, \varphi_{M-1}$ are given. Then we may evaluate Q by the fractional representation

$$Q = L \cdot \prod_{i=1}^{M-1} \varphi_i \cdot \left(\varphi_M \cdot m + \tau \cdot \left(\varphi_M \cdot h - \left(\sum_{i=1}^{M-1} \frac{1}{\varphi_i} + 1 \right) \right) \right) + \lambda \cdot h.$$

This costs $M+4$ multiplications, one subtraction, and $M+1$ additions, and hence the arithmetic complexities are $|Q_M| = M+4$, $|Q_S| = 1$, $|Q_A| = M+1$. Now, to attribute the inverses in formula (11), we increase the multiplicative complexity by $3 \cdot (M-1)$, which represents the fractional cost of the batch inversion⁵ of $\varphi_1, \dots, \varphi_{M-1}$. This yields the following equivalent complexities

$$|Q_M| = 4 \cdot M + 1, |Q_S| = 1, |Q_A| = M + 1,$$

which we may plug into formula (11).

Therefore the prover cost of Protocol 2 is as follows: Given the values of f_1, \dots, f_M and t over H , computing $\varphi_1 = x + f_1, \dots, \varphi_M = x + f_M$, $\tau = x + t$ costs $|H| \cdot (M+1) \cdot \mathbf{A}$, and their reciprocals $\frac{1}{\varphi_1}, \dots, \frac{1}{\varphi_M}$, $\frac{1}{\tau}$ are obtained within $3 \cdot |H| \cdot (M+1) \cdot \mathbf{M}$, using batch inversion. With these reciprocals we obtain the values for

$$h = \sum_{i=1}^M \frac{1}{\varphi_i} - \frac{m}{\tau}$$

⁵Batch, or Montgomery inversion, of a sequence $(a_i)_{i=1}^N$ computes the inverses of a_i^{-1} by recursively computing the cumulative products $p_i = a_1 \cdot \dots \cdot a_n$, $i = 0, \dots, n$, then calculating their inverses $q_i = \frac{1}{p_i}$ in a reverse manner starting with $q_n = \frac{1}{p_n}$, and putting $q_{i-1} = q_i \cdot a_i$, where i goes from n down to 1. The inverses are then derived via $a_i^{-1} = p_{i-1} \cdot q_i$, where $p_0 := 1$. The overall cost of the batch inversion is $3 \cdot (N-1)$ multiplications and a single inversion.

Table 1: Benchmark of Halo2’s Pippenger multi-scalar multiplication in the Pallas curve, varying the number N of scalars. The benchmarks were done on a AMD Ryzen 7 PRO 4750U, 32GB RAM DDR4, restricting to a single core.

$\log N$	Pippenger of size N	$2^8 \cdot N$ field mult.	equivalent field mult.
12	46.010 ms	21.11 ms	$N \cdot 557 \cdot M$
14	153.67 ms	84.78 ms	$N \cdot 464 \cdot M$
16	522.13 ms	169.70 ms	$N \cdot 394 \cdot M$
18	1.869 ms	679.27 ms	$N \cdot 351 \cdot M$

by $|H| \cdot (1 \cdot S + (M - 1) \cdot A)$. By the remark following Lemma 1, the values for $L_H(\vec{X}, \vec{y})$ over H are obtained within $|H| \cdot (M + A)$ operations. Hence the total cost of the preparation phase is

$$|H| \cdot ((3 \cdot M + 4) \cdot M + 1 \cdot S + (2 \cdot M + 1) \cdot A).$$

According to (11) the sumcheck costs

$$|H| \cdot \left(1 - \frac{1}{|H|}\right) \cdot ((5 \cdot M^2 + 24 \cdot M + 16) \cdot M + (2 \cdot M + 8) \cdot S + (2 \cdot M^2 + 13 \cdot M + 20) \cdot A).$$

However, as we may reuse the reciprocals of $\varphi_1, \dots, \varphi_{M-1}$ in the first step of the sumcheck, we correct the sumcheck cost by subtracting $|H| \cdot (3 \cdot (M - 1))$. Neglecting the $1/|H|$ -term, the overall cost of the prover is

$$|H| \cdot ((5 \cdot M^2 + 24 \cdot M + 23) \cdot M + (2 \cdot M + 9) \cdot S + (2 \cdot M^2 + 15 \cdot M + 21) \cdot A), \quad (22)$$

whereas it provides only two H -sized oracles. The cost is $\mathcal{O}(|H|)$ but depends quadratically in M the number of columns to be looked up. This quadratic occurrence is due to the fact that both, the number of function as well as the degree of Q grow linearly in M .

3.4 The optimal batch size

To determine the optimal number of columns M for Protocol 2, we compare its prover costs with those of the reference protocol from Section ???. For taking the oracle costs into account, we rely on the benchmarks from Table 1 which measure the equivalent number of field multiplication for a multi-scalar multiplication in an ordinary-sized elliptic curve with 128 bit security.

Based on these numbers, we were looking for M where the ratio of the overall number of field multiplications of the two protocols,

$$r = M(\text{Section } ??)/M(\text{Protocol 2}), \quad (23)$$

is maximal. These optimal batch sizes M , and their maximum ratios are given in Table 2.

As a lower bound for elliptic curve based schemes, we take the number of field operations for x^5 -Poseidon with rate $r = 2$, capacity $c = 1$, $R_F = 8$ full rounds and $R_p = 57$ partial rounds. (These are the parameters from [?] for a security level of 128 bits over a 255 bit large base field.) Each permutation, which processes $r = 2$ many elements costs $604 \cdot M + 591 \cdot A$, using the optimized evaluation strategy from Appendix B in [?]. Hence computing the hash of two functions over H costs

$$|H| \cdot (604 \cdot M + 591 \cdot A).$$

Therefore a M single-column lookups cost

$$|H| \cdot ((604 \cdot M + 42) \cdot M + 10 \cdot S + (591 \cdot M + 30) \cdot A), \quad (24)$$

Table 2: The number of columns M where the advantage of Protocol 2 over the reference variant from Section ?? is maximal. The numbers are based on the operation counts (22) and (40) for the oracle prover, and the benchmarks for a multi-scalar multiplication, Table 1.

$\log H $	12	14	16	18
M	13	12	11	10
r as in (23)	3.8	3.4	3.2	3.0

and the batch-column lookup for M columns

$$|H| \cdot ((5 \cdot M^2 + 19 \cdot M + 1226) \cdot \mathbf{M} + (2 \cdot M + 8) \cdot \mathbf{S} + (2 \cdot M^2 + 13 \cdot M + 621) \cdot \mathbf{A}). \quad (25)$$

Comparing (24) and (25) yields a maximum speedup by a factor of 5 for $M = 11$, more than double as fast for $3 \leq M \leq 58$, and a break-even point at about $M = 124$ columns.

3.5 Generalizations

Protocol 2 is easily generalized to functions with multilinear values,

$$\begin{aligned} t(\vec{x}) &= \sum_{(j_1, \dots, j_k) \in \{0,1\}^k} t_{j_1, \dots, j_k}(\vec{x}) \cdot Y_1^{i_1} \cdots Y_k^{j_k}, \\ f_i(\vec{x}) &= \sum_{(j_1, \dots, j_k) \in \{0,1\}^k} f_{i, j_1, \dots, j_k}(\vec{x}) \cdot Y_1^{i_1} \cdots Y_k^{j_k}, \end{aligned}$$

$i = 1, \dots, M$, without changing the soundness error bound from Theorem 4. As $F[X, Y_1, \dots, Y_k]$ is a unique factorization domain, and polynomials of the form $X - \sum_{(i_1, \dots, i_k) \in \{0,1\}^k} c_{i_1, \dots, i_k} \cdot Y_1^{i_1} \cdots Y_k^{i_k}$ are irreducible, we may apply Lemma 5 to see that $\bigcup_{i=1}^M \{f_i(\vec{x})\}_{\vec{x} \in H} \subseteq \{t(\vec{x})\}_{\vec{x} \in H}$ as sets in the rational function field $F(X, Y_1, \dots, Y_k)$, if and only if there exists a function $m : H \rightarrow F$ such that

$$\sum_{\vec{x} \in H} \sum_{i=1}^M \frac{1}{X + f_i(\vec{x})(\vec{Y})} = \sum_{\vec{x} \in H} \frac{m(\vec{x})}{X + t(\vec{x})(\vec{Y})}. \quad (26)$$

The only change to Protocol 2 is that the verifier now samples x from F and $\vec{y} = (y_1, \dots, y_k)$ from F^k , and continues with $x - f_i(\vec{x})$ and $x - t(\vec{x})$ replaced by $x - f_i(\vec{x})(\vec{y})$ and $x - t(\vec{x})(\vec{y})$.

4 Lookups based on the Hyperplonk shift

In this section we sketch batch-column lookups based on the Plookup strategy and the time shift from Hyperplonk+ [CBBZ22]. We give detailed operation counts and compare with the lookup argument from Section 3.

One of the main contributions of [CBBZ22] is the introduction of an (almost) transitive time shift $T : H \rightarrow H$ which translates to multilinear extensions in a tame manner. The shift is derived from the multiplication by a primitive root in $GF(2^n)$,

$$T(x_1, \dots, x_n) = \frac{1+x_n}{2} \cdot (1, x_1, \dots, x_{n-1}) + \frac{1-x_n}{2} \cdot (-1, (-1)^{1-c_1} \cdot x_1, \dots, (-1)^{1-c_{n-1}} \cdot x_{n-1}),$$

where the $c_i \in \{0, 1\}$ are the coefficients of a primitive polynomial $1 + \sum_{i=1}^{n-1} c_i \cdot X^i + X^n$ over $GF(2)$. The time shift acts transitively on the punctuated hypercube $H' = H \setminus \{\vec{1}\}$ (as a group automorphism it has $\vec{1}$ as a fixed point), and more importantly, evaluations of a shifted function $f(T(\vec{x}))$ can be simulated from two evaluations of f by

$$f(T(x_1, \dots, x_n)) = \frac{1+x_n}{2} \cdot f(1, x_1, \dots, x_{n-1}) + \frac{1-x_n}{2} \cdot f(-1, (-1)^{1-c_1} \cdot x_1, \dots, (-1)^{1-c_{n-1}} \cdot x_{n-1}). \quad (27)$$

Using the time shift T allows for the same strategy for the univariate Plookup argument. The argument is based on the fact that, given two sequences of field elements $(a_i)_{i=0}^{N-1}$ and $(t_i)_{i=0}^{N-1}$, then $\{a_i : i = 0, \dots, N-1\} \subseteq \{t_i : i = 0, \dots, N-1\}$ as sets, if and only if there exists a sequence $(s_i)_{i=0}^{2N-1}$ of double the size, which satisfies the identity

$$\prod_{i=0}^{N-1} (X + s_i + s_{i+1 \bmod N} \cdot Y) = \prod_{i=0}^{N-1} (X + a_i + a_i \cdot Y) \cdot (X + t_i + t_{i+1 \bmod N} \cdot Y). \quad (28)$$

(The sequence $(s_i)_{i=0}^{2N-1}$ is the concatenation of the (a_i) and (t_i) , ordered by value in the same way as in t .) We again discuss two approaches⁶ for dealing with the grand product obtained from (28). The main protocol, see Section 4.1, applies a batched grand product argument over H , independent of M the number of columns, leading to a $\mathcal{O}(M^2)$ prover for similar reasons as Protocol 2 does. The second one proves the grand product argument over an extended hypercube \bar{H} of size $M \cdot |H|$, which leads to $\mathcal{O}(M)$ oracle costs, but has bounded algebraic degree. We use this protocol as a reference to determine the optimal batch size.

4.1 Multivariate Plookup

Let $t : H' \rightarrow F$ be the lookup table, and $f_i : H' \rightarrow F$, $i = 1, \dots, M$, the functions subject to the lookup. Although the functions are defined over the punctuated hypercube H' , we assume arbitrary values at $\vec{1}$. (These will be ignored by the lookup argument.) The prover provides the ordered merge of the f_i together with t via additional functions $s_i : H' \rightarrow F$, $i = 1, \dots, M+1$, subject to the lookup identity

$$\begin{aligned} \prod_{\vec{x} \in H'} \prod_{i=1}^M (X + s_i(\vec{x}) + s_{i+1}(\vec{x}) \cdot Y) \cdot (X + s_{M+1}(\vec{x}) + s_1(T(\vec{x})) \cdot Y) \\ = \prod_{\vec{x} \in H'} \prod_{i=1}^M (X + f_i(\vec{x}) + f_i(\vec{x}) \cdot Y) \cdot (X + t(\vec{x}) + t(T(\vec{x})) \cdot Y). \end{aligned}$$

⁶We point out that the presented strategies slightly differ from the one in [CBBZ22], which uses the more expensive grand product argument from [SL20].

The identity is reduced to a grand product over H' by random samples $\alpha, \beta \leftarrow_s F$ for X and Y , yielding

$$\prod_{\vec{x} \in H'} h(\vec{x}) = 1,$$

where

$$h(\vec{x}) = \frac{\alpha + s_{M+1}(\vec{x}) + s_1(T(\vec{x})) \cdot \beta}{\alpha + t(\vec{x}) + t(T(\vec{x})) \cdot \beta} \cdot \prod_{i=1}^M \frac{(\alpha + s_i(\vec{x}) + s_{i+1}(\vec{x}) \cdot \beta)}{(\alpha + f_i(\vec{x}) + f_i(\vec{x}) \cdot \beta)}.$$

The prover computes the cumulative products of the values $h(\vec{x})$ along the orbit of the time shift T on H' , starting with $\phi(-\vec{1}) = 1$, and setting

$$\phi(T^k(-\vec{1})) = \phi(T^{k-1}(-\vec{1})) \cdot h(T^{k-1}(-\vec{1})),$$

for $k = 1, \dots, |H| - 1$. At the remaining point $\vec{x} = \vec{1}$ outside H' , the prover sets $\phi(\vec{x})$ to zero. Correctness of the grand product is proven by the constraint on its initial value $\phi(-\vec{1}) = 1$, and the domain identity

$$\phi(T(\vec{x})) \cdot \tau(\vec{x}) \cdot \prod_{i=1}^M \varphi_i(\vec{x}) - \phi(\vec{x}) \cdot \sigma_{M+1}(\vec{x}) \cdot \prod_{i=1}^M \sigma_i(\vec{x}) = 0, \quad (29)$$

for all $\vec{x} \in H$, where $\tau(\vec{x}) = \alpha + t(\vec{x}) + \beta \cdot t(T(\vec{x}))$ and

$$\begin{aligned} \varphi_i(\vec{x}) &= \alpha + (1 + \beta) \cdot f_i(\vec{x}), \\ \sigma_i(\vec{x}) &= \alpha + s_i(\vec{x}) + \beta \cdot s_{i+1}(\vec{x}), \end{aligned}$$

for $i = 1, \dots, M$, except $\sigma_{M+1}(\vec{x}) = \alpha + s_{M+1}(\vec{x}) + \beta \cdot s_1(T(\vec{x}))$. As in Protocol 2, both constraints on ϕ , the initial value condition and the domain identity, are reduced to sumchecks over H by help of the Lagrange polynomials $L_H(\cdot, -\vec{1})$ and $L_H(\cdot, \vec{y})$, where $\vec{y} \leftarrow_s F^n$, and then combined into a single one by a batching randomness $\lambda \leftarrow_s F$. The resulting overall sumcheck is

$$\sum_{\vec{x} \in H} Q(L_H(\vec{x}, \vec{y}), L_H(\vec{x}, -\vec{1}), \tau(\vec{x}), \varphi_1(\vec{x}), \dots, \varphi_M(\vec{x}), \sigma_1(\vec{x}), \dots, \sigma_{M+1}(\vec{x}), \phi(\vec{x}), \phi(T(\vec{x}))) = 0,$$

where

$$Q(L_H, L, \tau, \varphi_1, \dots, \varphi_M, \sigma_1, \dots, \sigma_{M+1}, \phi, \phi_T) = L_H \cdot \left(\phi_T \cdot \tau \cdot \prod_{i=1}^M \varphi_i - \phi \cdot \prod_{i=1}^{M+1} \sigma_i \right) + \lambda \cdot L \cdot (\phi - 1). \quad (30)$$

Note that Q has absolute degree $d = M + 3$, which is the same as in Protocol 2, but about the double of variables, $\nu = 2 \cdot M + 6$. Its arithmetic complexities are $|Q_M| = 2 \cdot (M + 1) + 3$, $|Q_S| = 2$, $|Q_A| = 1$.

4.2 A reference variant for Section 4.1

Suppose that we have a table function $t : H' \rightarrow F$ defined over the punctuated hypercube $H' = H \setminus \{\vec{1}\}$, and M column functions $f_1, \dots, f_M : H \rightarrow F$ defined on the *entire* hypercube. We assume that $M + 1 = 2^m$ for an integer $m \geq 1$, so that we can index the column functions as $f_{\vec{z}}$, with \vec{z} from $\{\pm 1\}^m \setminus \{\vec{1}\}$. The ordered concatenation of the column functions and the table function has now $2^m \cdot |H| - 1$ entries, which we arrange in a single function s over $\bar{H} \setminus \{\vec{1}\}$ along the orbit of the time shift $T_{\bar{H}}$ on \bar{H} , where $\bar{H} = \{\pm 1\}^m \times H$. The value of s at $\vec{1}$ can be set arbitrary (it will be ignored by the lookup argument). Using the patched function

$$f(\vec{y}, \vec{x}) = \sum_{\vec{z} \in \{\pm 1\}^m \setminus \{\vec{1}\}} L_m(\vec{y}, \vec{z}) \cdot f_{\vec{z}}(\vec{x}),$$

over \bar{H} , where L_m is the Lagrange kernel for $\{\pm 1\}^m$, the lookup identity at random $(\alpha, \beta) \leftarrow_{\$} F^2$ reads now as

$$\prod_{(\vec{y}, \vec{x}) \in \bar{H}'} (\alpha + s(\vec{y}, \vec{x}) + \beta \cdot s(T_{\bar{H}}(\vec{y}, \vec{x}))) = \prod_{(\vec{y}, \vec{x}) \in \bar{H}'} (\alpha + (1 + \beta) \cdot f(\vec{y}, \vec{x}) + L_m(\vec{y}, \vec{1}) \cdot (t(\vec{x}) + \beta \cdot t(T_H(\vec{x}))),$$

where $\bar{H}' = \bar{H} \setminus \{\vec{1}\}$. (Notice that the α -term for t on the right hand side is combined with the α -term for f .) In other words the prover needs to show that

$$\prod_{(\vec{y}, \vec{x}) \in \bar{H}'} \frac{\varphi(\vec{y}, \vec{x})}{\sigma(\vec{y}, \vec{x})} = 1,$$

where

$$\begin{aligned} \varphi(\vec{y}, \vec{x}) &= \alpha + (1 + \beta) \cdot f(\vec{y}, \vec{x}) + L_m(\vec{y}, \vec{1}) \cdot (t(\vec{x}) + \beta \cdot t(T_H(\vec{x}))), \\ \sigma(\vec{y}, \vec{x}) &= \alpha + s(\vec{y}, \vec{x}) + \beta \cdot s(T_{\bar{H}}(\vec{y}, \vec{x})). \end{aligned}$$

As before, the prover provides a function ϕ over \bar{H}' for the cumulative products of the values of $h(\vec{x}) = \frac{\varphi(\vec{y}, \vec{x})}{\sigma(\vec{y}, \vec{x})}$ along the orbit of $T_{\bar{H}}$, starting with $\phi(-\vec{1}) = 1$, and setting

$$\phi(T_{\bar{H}}^k(-\vec{1})) = \phi(T_{\bar{H}}^{k-1}(-\vec{1})) \cdot h(T_{\bar{H}}^{k-1}(-\vec{1})),$$

for $k = 1, \dots, |\bar{H}| - 2$. At the remaining point $\vec{1}$ we set $\phi(\vec{1}) = 0$. Correctness of ϕ is ensured by $\phi(-\vec{1}) = 1$, and the domain identity

$$\phi(T_{\bar{H}}(\vec{y}, \vec{x})) \cdot \varphi(\vec{y}, \vec{x}) - \phi(\vec{y}, \vec{x}) \cdot \sigma(\vec{y}, \vec{x}) = 0$$

over \bar{H} . Applying the Lagrange polynomials $L_{\bar{H}}(\cdot, \vec{z})$, $\vec{z} \leftarrow_{\$} F^{m+n}$, and $L_{\bar{H}}(\cdot, -\vec{1})$, the constraints on ϕ are reduced to sumchecks over \bar{H} , which are then combined into a single one using a random $\lambda \leftarrow_{\$} F$. The overall sumcheck is

$$\sum_{(\vec{y}, \vec{x}) \in \bar{H}} Q(L_{\bar{H}}((\vec{y}, \vec{x}), \vec{z}), L_{\bar{H}}((\vec{y}, \vec{x}), -\vec{1}), \varphi(\vec{y}, \vec{x}), \sigma(\vec{y}, \vec{x}), \phi(\vec{y}, \vec{x}), \phi(T_{\bar{H}}(\vec{y}, \vec{x}))) = 0,$$

where

$$Q(L_H, L, \varphi, \sigma, \phi, \phi_T) = L_H \cdot (\phi_T \cdot \varphi - \phi \cdot \sigma) + \lambda \cdot L \cdot (\phi - 1). \quad (31)$$

The polynomial Q has $\nu = 6$ variables, absolute degree $d = 3$, and its arithmetic complexities are $|Q_M| = 5$, $|Q_S| = 2$, $|Q_A| = 1$.

4.3 Computational cost

The prover cost for the Protocol from Section 4.1 is as follows. Computing the values for τ and all φ_i, σ_i over H consumes overall

$$|H| \cdot (2 \cdot (M + 1) \cdot M + (3M + 4) \cdot A),$$

the quotient $h(\vec{x})$ is obtained within $|H| \cdot (2 \cdot M + 4) \cdot M$, using batch inversion. From these values $\phi(x)$ over H is derived by another $|H| \cdot M$. The domain evaluation for $L_H(\vec{X}, \vec{y})$ is obtained within $|H| \cdot (M + A)$ operations, and the sumcheck costs

$$|H| \cdot \left(1 - \frac{1}{|H|}\right) \cdot ((4 \cdot M^2 + 25 \cdot M + 38) \cdot M + (6 \cdot M + 20) \cdot S + (2 \cdot M^2 + 14 \cdot M + 26) \cdot A).$$

Neglecting the $1/|H|$ -term, the overall cost of the prover is

$$|H| \cdot ((4 \cdot M^2 + 29 \cdot M + 46) \cdot M + (6 \cdot M + 20) \cdot S + (2 \cdot M^2 + 17 \cdot M + 31) \cdot A). \quad (32)$$

This is comparable with (22), but does not take into account that the prover needs to supply a total of $M + 2$ functions over H instead of two.

The cost of our reference protocol from Section 4.2 is as follows. Computing the values of the patched function φ over \bar{H} takes $|\bar{H}| \cdot (1 \cdot M + 2 \cdot A)$, and the same number of operations are needed for σ . The quotient h over \bar{H} is obtained by batch inversion, consuming

$$|\bar{H}| \cdot (3 + 1) \cdot M,$$

and the values of ϕ are computed within another $|\bar{H}|$ multiplications. Last but not least, computing $L_{\bar{H}}(\cdot, \vec{z})$ over \bar{H} costs $|\bar{H}| \cdot (M + A)$, and the sumcheck

$$|\bar{H}| \cdot \left(1 + \frac{1}{|\bar{H}|}\right) \cdot (38 \cdot M + 20 \cdot S + 26 \cdot A).$$

Neglecting the $\frac{1}{|\bar{H}|}$ -term, the overall cost of the prover is

$$(M + 1) \cdot |H| \cdot (45 \cdot M + 20 \cdot S + 30 \cdot A), \quad (33)$$

whereas it needs to provide the functions s and ϕ over \bar{H} , which amounts the equivalent of $2 \cdot (M + 1)$ functions over $|H|$.

Based on our benchmark-backed equivalent number of field operations for a multi-scalar multiplications, Table 1, we obtain the optimal cost ratio

$$r = \frac{M(\text{Section 4.1})}{M(\text{Section 4.2})} \approx 1.7$$

throughout at a two-adic batch size of $M = 8$, considering domain sizes $|H|$ between $\log |H| = 12$ and $\log |H| = 18$.

4.4 Comparison with logarithmic derivative lookups

The significant advantage of the logarithmic derivative lookup over the one from the current section is the lower oracle costs. While the Plookup strategy demands $(M + 2)$ oracles, the logarithmic derivative only uses 2, while having comparable arithmetic costs. We compare the two strategies by their computational cost per column, when running them at optimal batch size M .

5 Acknowledgements

The author would like to thank Rayan Matovu and Morgan Thomas for giving me the space and time to dwell on batch-column lookups. Special thanks to Marcin Bugaj for helping out with the Pippenger benchmarks. Furthermore, I would like to thank Ariel Gabizon for his feedback and useful discussions.

Table 3: The estimated performance advantage of logarithmic derivative lookups (Protocol 2) over the multivariate Plookup strategy (Section 4.1), as the ratio of their number of field multiplications per column. The numbers are based on the equivalent number of field multiplications for a multi-scalar multiplication over the Pallas curve, given in Table 1.

$\log H $	Number of field mult. per column		ratio
	plookup strat.	log. derivative	
12	764 ($M = 8$)	177 ($M = 13$)	4.3
14	646 ($M = 8$)	163 ($M = 12$)	4.0
16	559 ($M = 8$)	153 ($M = 11$)	3.7
18	506 ($M = 8$)	146 ($M = 10$)	3.5

Table 4: The estimated performance advantage of logarithmic derivative lookups over the ones from this section, as the ratio of their number of field multiplications $r = \mathsf{M}(\text{Plookup})/\mathsf{M}(\log D)$. The numbers are based on the equivalent number of field multiplications from Table 1. For hypercube sizes $|H|$ ranging from 2^{12} – 2^{18} we describe the maximum ratio r_{\max} over the number of columns M , as well as the ranges for M over which r is larger than 2 and 3, respectively. The minimum ratio is throughout $r = 1.5$ and obtained in the single-column setting $M = 1$.

$\log H $	$r \geq 3$	$r \geq 2$	r_{\max}
12	$M \in [5, 41]$	$M \in [3, 87]$	4.1 (at $M = 15$)
14	$M \in [6, 32]$	$M \in [3, 71]$	3.8 (at $M = 13$)
16	$M \in [6, 26]$	$M \in [3, 59]$	3.5 (at $M = 12$)
18	$M \in [6, 21]$	$M \in [3, 52]$	3.2 (at $M = 12$)

References

- [BCC⁺16] Jonathan Bootle, Andrea Cerulli, Pyrros Chaidos, Jens Groth, and Christophe Petit. Efficient zero-knowledge arguments for arithmetic circuits in the discrete log setting. In M. Fischlin and J.S. Coron, editors, *EUROCRYPT 2016*, volume 9666 of *LNCS*. Springer, 2016. Full paper: <https://eprint.iacr.org/2016/263>.
- [BCG⁺18] Jonathan Bootle, Andrea Cerulli, Jens Groth, Sune Jakobsen, and Mary Maller. Arya: Nearly linear-time zero-knowledge proofs for correct program execution. In *ASIACRYPT 2018*, volume 11272 of *LNCS*, 2018. Full paper: <https://eprint.iacr.org/2018/380>.
- [BFS20] Benedikt Bünz, Ben Fisch, and Alan Szepieniec. Transparent SNARKs from DARK compilers. In *EUROCRYPT 2020*, 2020. Full paper: <https://eprint.iacr.org/2019/1229>.
- [BG12] Stephanie Bayer and Jens Groth. Efficient zero-knowledge argument for correctness of a shuffle. In *EUROCRYPT 2012*, volume 7237 of *LNCS*. Springer, 2012.
- [BSBHR18] Eli Ben-Sasson, Iddo Bentov, Yinon Horesh, and Michael Riabzev. Scalable, transparent, and post-quantum secure computational integrity. In *IACR ePrint Archive 2018/046*, 2018. <https://eprint.iacr.org/2018/046>.
- [BSCS16] Eli Ben-Sasson, Alessandro Chiesa, and Nicholas Spooner. Interactive oracle proofs. In *TCC 2016*, pages 31–60, 2016.

- [CBBZ22] Binyi Chen, Benedikt Bünz, Dan Boneh, and Zhenfei Zhang. Hyperplonk: PLONK with a linear-time prover and high-degree costumed gates. In *IACR ePrint Archive 2020/1355*, 2022. <https://eprint.iacr.org/2022/1355>.
- [Gab22] Ariel Gabizon. Multiset checks in PLONK and Plookup. hackmd.io, 2022. <https://hackmd.io/@arielg/ByFgSDA7D>.
- [GK22] Ariel Gabizon and Dmitry Khovratovich. flookup: Fractional decomposition-based lookups in quasi-linear time independent of the table size. In *IACR ePrint Archive 2022/1447*, 2022. <https://eprint.iacr.org/2022/1447>.
- [GW20] Ariel Gabizon and Zachary J. Williamson. Plookup: A simplified polynomial protocol for lookup tables. In *IACR ePrint Archive 2020/315*, 2020. <https://eprint.iacr.org/2020/315>.
- [GWC19] Ariel Gabizon, Zachary J. Williamson, and Oana Ciobotaru. PLONK: Permutations over Lagrange-bases for oecumenical non-interactive arguments of knowledge. In *IACR ePrint Archive 2019/953*, 2019. <https://eprint.iacr.org/2019/953>.
- [KZG10] Aniket Kate, Gregory M. Zaverucha, and Ian Goldberg. Constant-size commitments to polynomials and their applications. In Abe M., editor, *ASIACRYPT 2010*, volume 6477 of *LNCS*. Springer, 2010.
- [Lab] Avi Dessauer (Orbis Labs). tiny-ram-halo2. <https://github.com/Orbis-Tertius/tiny-ram-halo2>.
- [LFKN92] Carsten Lund, Lance Fortnow, Howard Karloff, and Noam Nisan. Algebraic methods for interactive proof systems. In *Journal of the Association for Computing Machinery*, volume 39, pages 859–868, 1992. Full paper: <https://eprint.iacr.org/2017/1066>.
- [PK22] Jim Posen and Assimakis A. Kattis. Caulk+: Table-independent lookup arguments. In *IACR ePrint Archive 2022/957*, 2022. <https://eprint.iacr.org/2022/957>.
- [Pol] Polygon Zero Team: plonky2. <https://github.com/mir-protocol/plonky2>.
- [PST13] Charalampos Papamanthou, Elaine Shi, and Roberto Tamassia. Signatures of correct computation. In *TCC 3*, volume 7785 of *LNCS*. Springer, 2013.
- [RDJH22] Gyumin Roh, Wei Dai, Maria Jabbour, and Andrew He. New lookup argument. zkResearch, 2022. <https://zkresearch.ch/t/new-lookup-argument/32>.
- [SL20] Srinath Setty and Jonathan Lee. Quarks: Quadruple-efficient transparent zkSNARKs. In *IACR ePrint Archive 2020/1275*, 2020. <https://eprint.iacr.org/2020/1275>.
- [Tha13] Justin Thaler. Time-optimal interactive proofs for circuit evaluation. In *CRYPTO 2013*, volume 8043 of *LNCS*, 2013. Full paper: <https://eprint.iacr.org/2013/351>.
- [ZBK⁺22] Arantxa Zapico, Vitalik Buterin, Dmitry Khovratovich, Mary Maller, Anca Nitulescu, and Mark Simkin. Caulk: Lookup arguments in sublinear time. In *IACR ePrint Archive 2022/621*, 2022. <https://eprint.iacr.org/2022/621>.

A Appendix: the univariate case

In this section we compare the univariate variant of Protocol 2 with the lookups from [GW20] and [GK22], generalized to batches of columns. As the “large number of columns” batching strategy from Section ?? again yields a break-even point beyond practical interest⁷, we skip a separate analysis of it. The protocols are formulate as Lagrange IOPs, even though the advantage over polynomial oracles is modest: For most polynomials the prover needs to compute their coefficients anyway.

We throughout assume that F is an FFT-friendly finite field, having a multiplicative subgroup $H = \{x \in F : x^n = 1\}$ of order n , and denote by g a generator of it. The Lagrange kernel of H is the symmetric bivariate polynomial

$$L_H(X, Y) = \frac{1}{n} \cdot \left(1 + \sum_{i=1}^{n-1} X^i \cdot Y^{n-i} \right) = \frac{1}{n} \cdot \frac{Y \cdot v_H(X) - X \cdot v_H(Y)}{X - Y},$$

where $v_H(X) = X^n - 1$ is the vanishing polynomial of H . Again, for $y \in H$ the polynomial $L_H(X, y)$ is the Lagrange polynomial which is equal to 1 at $x = y$, and zero elsewhere on H . In all the protocols, the prover is given $M \geq 1$ functions $f_i : H \rightarrow F$, $i = 0, \dots, M-1$, and wants to prove that their ranges are contained in that of a prediscrbed table $t : H \rightarrow F$, i.e.

$$\bigcup_{i=0}^{M-1} \{f_i(x)\}_{x \in H} \subseteq \{t(x)\}_{x \in H}.$$

For notational convenience, will use the same notation for the function $f : H \rightarrow F$ and its sequence of values $(f(g^k))_{k=0}^{n-1}$.

Our cost analysis is based on the following approximate formula: Assume that we have Lagrange representation of $\nu \geq 1$ polynomials $p_1(X), \dots, p_\nu(X)$ over H , where all $\deg p_i(X) < |H|$, and an ν -variate polynomial Q over F with absolute degree d and multiplicative complexity Q_M . Then the number of field multiplications for computing the Lagrange representation of the quotient polynomial $q(X)$ in

$$Q(p_1(X), \dots, p_\nu(X)) = 0 \pmod{v_H(X)},$$

over the “multiplication domain” H' of size $d \cdot |H|$ (except the trivial coset H) is

$$\nu \cdot \left(\frac{|H|}{2} \cdot \log |H| + \frac{d \cdot |H|}{2} \cdot \log(d \cdot |H|) \right) + d \cdot |H| \cdot (Q_M + 1) \quad (34)$$

The first term corresponds to the computation of the coefficients of $p_1(X), \dots, p_\nu(X)$, plus their values over the multiplication domain, and the second term is for the pointwise evaluation of Q over that domain, plus the multiplication with the precomputed inverses of the vanishing polynomial for obtaining the values of $q(X)$. Theses values over the $d-1$ (non-trivial) cosets of H , enumerated as $a^i \cdot H$, $i = 1, \dots, d-1$, where a is any element of $H' \setminus H$, define the Lagrange representations of “component polynomials” $q_0(X), \dots, q_M(X)$ of degree $\deg p_i(X) < |H|$ in the decomposition

$$q(X) = \frac{1}{d} \cdot \sum_{i=1}^{d-1} \frac{v_{H'}(a^{-i} \cdot X)}{v_H(a^{-i} \cdot X)} \cdot q_i(a^{-i} \cdot X). \quad (35)$$

We point out that when embedding lookups in other IOPs (for example Plonk), so that their overall identities share the same quotient polynomial, then the fractional cost for the lookup is less.

⁷For example, at 2^{12} constraints we obtain a break-even point at about $M = 40$ columns.

A.1 The logarithmic derivate approach

The batch-column argument

The lookup condition is equivalent to that the rational function

$$h(x) = \sum_{i=0}^{M-1} \frac{1}{\alpha + f_i(x)} - \frac{m(x)}{\alpha + t(x)}, \quad (36)$$

sums up to zero over H , where $\alpha \leftarrow_{\$} F$ is the first verifier challenge. For this the prover provides the Lagrange representation of $\phi(X)$ subject to the domain identity

$$\phi(g \cdot x) - \phi(x) = \sum_{i=0}^{M-1} \frac{1}{\alpha + f_i(x)} - \frac{m(x)}{x + t(x)},$$

or

$$\tau(X) \cdot \prod_{i=0}^{M-1} \varphi_i(X) \cdot (\phi(g \cdot X) - \phi(X)) = \tau(X) \cdot \prod_{i=0}^{M-1} \varphi_i(X) \cdot \left(\sum_{i=0}^{M-1} \frac{1}{\varphi_i(X)} - \frac{m(X)}{\tau(X)} \right) \mod v_H(X), \quad (37)$$

where the right hand side is a polynomial, and

$$\begin{aligned} \varphi_i(X) &= \alpha + f_i(X), \quad i = 0, \dots, M-1, \\ \tau(X) &= \alpha + t(X). \end{aligned}$$

The overall identity (37) is of the form

$$Q(\tau(X), \varphi_0(X), \dots, \varphi_{M-1}(X), \phi(X), \phi(g \cdot X)) = 0 \mod v_H(X),$$

where Q has $\nu = M + 3$ variables and absolute degree $d = M + 2$. The prover computes the Lagrange representation of the overall quotient polynomial over the multiplication domain H' of size $(M + 1) \cdot |H|$, except H , and sends the Lagrange representations of the component polynomials $q_1(X), \dots, q_M(X)$ as defined by (35) to the verifier. The verifier checks the overall identity at a random $\beta \leftarrow_{\$} F$ by querying the oracles at the needed points.

Cost analysis

As in Section 3, we use the fractional representation from (37) for computing domain evaluations of Q . Using batch inversion, the equivalent multiplicative complexity for computing Q is

$$Q_M = M + 1 + 3 \cdot (M + 1) + 2 = 4 \cdot M + 6,$$

taking into account the proportional cost of batch inversion as 3 multiplications.

The prover cost, simplified to the number of field multiplications, is as follows. Computing the values of $\varphi_0, \dots, \varphi_{M-1}$ and τ over H does not cost a single field multiplication, their inverses are obtained within

$$(M + 1) \cdot 3 \cdot |H|$$

multiplications, from which $h(x)$ over H and the values of the sumcheck polynomial ϕ are obtained by another overall $2 \cdot |H|$ multiplications. Since the values of the shifted function $\phi(g \cdot X)$ are directly obtained

from those of ϕ , we may apply (34) for $\nu = M + 2$, and computing the Lagrange representation of the overall quotient takes

$$(M + 2) \cdot \left(\frac{|H|}{2} \cdot \log |H| + \frac{(M + 2) \cdot |H|}{2} \cdot \log((M + 2) \cdot |H|) \right) + (M + 2) \cdot |H| \cdot (Q_M + 1)$$

multiplications. Overall, the oracle prover demands

$$|H| \cdot \left(\frac{(M + 2)^2}{2} \cdot (\log |H| + \log(M + 2)) + (M + 2) \cdot (4 \cdot M + 6) + \frac{M + 2}{2} \cdot \log |H| + 4 \cdot M + 7 \right) \quad (38)$$

field multiplications, and it provides the equivalent of $M + 3$ functions of size $|H|$, i.e. the values of m , ϕ and the component polynomials q_1, \dots, q_{M+1} over H .

The reference variant

The prover provides $M + 1$ helper functions, one for each reciprocal $h_i(x) = \frac{1}{\alpha + f_i(x)}$, $i = 1, \dots, M$, and another one for $h_t(x) = \frac{1}{\alpha + t(x)}$. Correctness of the helper functions is enforced by the domain identities

$$h_i(x) \cdot (\alpha + f_i(x)) - 1 = 0,$$

for $i = 1, \dots, M$, and

$$h_t(x) \cdot (\alpha + t(x)) - 1 = 0,$$

and these are altogether subject to the sumcheck

$$\sum_{i=1}^M h_i(x) - h_t(x) \cdot m(x) = \phi(g \cdot x) - \phi(x).$$

The domain identities are again turned into a sumchecks by applying the Lagrange kernel $L_H(\cdot, \vec{z})$, where $\vec{z} \leftarrow_{\$} F^n$. Combining the sumchecks using the powers of a random $\lambda \leftarrow_{\$} F$ leads to the overall sumcheck

$$\sum_{\vec{x} \in H} Q(L_H(\vec{x}, \vec{z}), \varphi_1(\vec{x}), \dots, \varphi_M(\vec{x}), \tau(\vec{x}), h_1(\vec{x}), \dots, h_M(\vec{x}), h_t(\vec{x}), m(\vec{x})) = 0,$$

with

$$Q(L, \varphi_1, \dots, \varphi_M, \tau, h_1, \dots, h_M, h_t) = \sum_{i=1}^M h_i - m \cdot h_t + L \cdot \left(\sum_{i=1}^M \lambda^i \cdot (h_i \cdot \varphi_i - 1) + \lambda^{M+1} \cdot (h_t \cdot \tau - 1) \right). \quad (39)$$

In this variant, the prover provides an overall equivalent of $M + 2$ oracles of size $|H|$. The sumcheck polynomial Q has $\nu = 2 \cdot (M + 2)$ variables but its absolute degree $d = 3$ is independent of the number of columns. The arithmetic complexities of Q are $|Q_M| = 2 \cdot (M + 2)$, $|Q_S| = M + 2$, $|Q_A| = 2 \cdot M$.

Computational cost

The prover cost of the protocol is as follows. Computing the values of h_1, \dots, h_M and h_t over H using batch inversion costs

$$(M + 1) \cdot |H| \cdot (3 \cdot M + A),$$

and the values for $L_H(\cdot, \vec{z})$ over H are determined in $|H| \cdot (M + A)$. According to (11) the overall sumcheck costs

$$|H| \cdot \left(1 - \frac{1}{|H|} \right) \cdot (14 \cdot (M + 2) \cdot M + 6 \cdot (M + 2)S + (7 \cdot M + 16) \cdot A).$$

Neglecting the $\frac{1}{M \cdot |H|}$ -term, the overall cost for the prover is

$$|H| \cdot ((17 \cdot M + 32) \cdot M + (6 \cdot M + 12) \cdot S + (15 \cdot M + 18) \cdot A), \quad (40)$$

but the prover needs to provide the oracles for one function over a domain of size $M \cdot |H|$, and one over H .

A.2 Plookup

We describe the [Gab22] variant of plookup [GW20], generalized to batch-lookups. Consider the concatenation of the sequences $(f_i(g^k))_{k=0}^{n-1}$, $i = 0, \dots, M-1$, and (t_k) , ordered as they occur in (t_k) :

$$\bar{s} = (\bar{s}_i)_{i=0}^{(M+1) \cdot n-1} = (\underbrace{t_0, \dots, t_0}_{1+m_0 \text{ times}}, \underbrace{t_1, \dots, t_1}_{1+m_1 \text{ times}}, \dots, \underbrace{t_{n-1}, \dots, t_{n-1}}_{1+m_{n-1} \text{ times}}),$$

and split it into $(M+1)$ sequences of length n ,

$$s_i = (\bar{s}_{i+k \cdot (M+1)})_{k=0}^{n-1},$$

with $i = 0, \dots, M$. We regard these sequences again as functions on H . Then

$$\bigcup_{i=0}^{M-1} \{s_i(x), s_{i+1}(x)\}_{x \in H} \cup \{(s_M(x), s_0(g \cdot x))\}_{x \in H} = \bigcup_{i=0}^{M-1} \{(f_i(x), f_i(x))\}_{x \in H} \cup \{(t(x), t(g \cdot x))\}_{x \in H} \quad (41)$$

as multisets. Moreover, this multiset equality is equivalent to $\bigcup_{i=0}^{M-1} \{f_i(x)\}_{x \in H} \subseteq \{t(x)\}_{x \in H}$. That is, the set inclusion holds if and only if there exists functions $s_0, \dots, s_M : H \rightarrow F$ satisfying (41), which is rephrased by the lookup identity

$$\begin{aligned} \prod_{x \in H} \prod_{i=0}^{M-1} (X + s_i(x) + s_{i+1}(x) \cdot Y) \cdot (X + s_M(x) + s_0(g \cdot x) \cdot Y) \\ = \prod_{x \in H} \prod_{i=0}^{M-1} (X + f_i(x) + f_i(x) \cdot Y) \cdot (X + t(x) + t(g \cdot x) \cdot Y). \end{aligned} \quad (42)$$

The lookup protocol is as follows. The prover sends s_0, \dots, s_M subject to (42) to the verifier, which returns random samples $\alpha, \beta \leftarrow F$ for X and Y , reducing the lookup identity to the grand product

$$\prod_{x \in H} \frac{\sigma_M(x)}{\tau(x)} \cdot \prod_{i=0}^{M-1} \frac{\sigma_i(x)}{\varphi_i(x)} = 1, \quad (43)$$

where

$$\begin{aligned} \tau(x) &= \alpha + t(x) + \beta \cdot t(g \cdot x), \\ \varphi_i(x) &= \alpha + (1 + \beta) \cdot f_i(x), \quad i = 0, \dots, M-1, \\ \sigma_i(x) &= \alpha + s_i(x) + \beta \cdot s_{i+1}(x), \quad i = 0, \dots, M-1, \\ \sigma_M(x) &= \alpha + s_M(x) + \beta \cdot s_0(g \cdot x). \end{aligned}$$

To prove (43), the prover provides the cumulative products polynomial ϕ for the function $h(x) = \frac{\sigma_M(x)}{\tau(x)} \cdot \prod_{i=0}^{M-1} \frac{\sigma_i(x)}{\varphi_i(x)}$, which is characterized by the domain identities

$$\phi(g \cdot x) \cdot \tau(x) \cdot \prod_{i=0}^{M-1} \varphi_i(x) - \phi(x) \cdot \prod_{i=0}^M \sigma_i(x) = 0,$$

and

$$L_H(x, 1) \cdot (\phi(x) - 1) = 0,$$

for all $x \in H$. The identities are combined into a single one using a random $\lambda \leftarrow_{\$} F$ from the verifier, yielding the overall identity

$$Q(L_H(X, 1), \varphi_0(X), \dots, \varphi_{M-1}(X), \sigma_0(X), \dots, \sigma_M(X), \tau(X), \phi(X), \phi(g \cdot X)) = 0 \pmod{v_H(X)}, \quad (44)$$

where

$$Q(L, \varphi_0, \dots, \varphi_{M-1}, \sigma_0, \dots, \sigma_M, \tau, \phi, \phi_g) = \phi_g \cdot \tau \cdot \prod_{i=0}^{M-1} \varphi_i - \phi \cdot \prod_{i=0}^M \sigma_i + \lambda \cdot L \cdot (\phi - 1). \quad (45)$$

The prover provides (the Lagrange representation) of the quotient components $q_1(X), \dots, q_{M+1}(X)$, and the verifier checks the overall identity (44) at a random $\gamma \leftarrow_{\$} F$, by querying the involved polynomials at the needed points.

Cost analysis

The polynomial Q in the overall identity 44 has $\nu = 2 \cdot (M + 1) + 3$ variables, absolute degree $d = M + 2$, and multiplicative complexity $Q_M = 2 \cdot (M + 1) + 2$. Note that the values of $\phi(g \cdot X)$ over H or any larger domain can be obtained by shifting, hence we can reduce the number of variables by one when using the cost formula (34). The values over H for all φ_i , σ_i , and τ (overall $2 \cdot (M + 1)$ functions) consume

$$|H| \cdot 2 \cdot (M + 1),$$

field multiplications, the rational function $h(x)$ over H is derived within

$$6 \cdot |H|,$$

using batch inversion, and computing the cumulative product function ϕ over H costs another $|H| \cdot M$. Assuming that the values of $L_H(x, 1)$ over the multiplication domain are precomputed, and that the values of $\phi(g \cdot X)$ over that domain are obtained by shifting, the cost of computing the overall quotient polynomial is

$$(2 \cdot M + 3) \cdot \left(\frac{|H|}{2} \cdot \log |H| + \frac{d \cdot |H|}{2} \cdot \log(d \cdot |H|) \right) + d \cdot |H| \cdot (Q_M + 1).$$

The overall cost of the oracle prover is

$$|H| \cdot \left(\frac{(2 \cdot M + 3) \cdot (M + 3)}{2} \cdot \log |H| + (M + 2) \cdot (2 \cdot M + 4) + \frac{(2 \cdot M + 3) \cdot (M + 2)}{2} \cdot \log d + 2 \cdot M + 9 \right), \quad (46)$$

while the oracle costs are equivalent to $2 \cdot (M + 1) + 1$ domain evaluations of size $|H|$, corresponding to s_0, \dots, s_M , ϕ , and the quotient components q_1, \dots, q_{M+1} .

A.3 The flookup strategy

Section 5 of [GK22] describes a polynomial IOP for lookups, which is almost identical to the logarithmic derivative approach. We present its generalization to batch-lookups. For showing that the ranges of witness function $f_i : H \rightarrow F$, $i = 0, \dots, M - 1$, are contained in the range of a table $t : H \rightarrow F$, the fractional logarithmic derivative identity is turned into the polynomial identity

$$\sum_{x \in H} \sum_{i=0}^{M-1} \frac{v_T(X)}{X + f_i(x)} = \sum_{x \in H} m(x) \cdot \frac{v_T(X)}{X + t(x)},$$

by multiplying with the precomputed table polynomial $v_T(X) = \prod_{x \in H} (X + t(x))$. Instead of the multiplicity function m , the prover explicitly provides the polynomial⁸

$$R_T(X) = \sum_{x \in H} m(x) \cdot \frac{v_T(X)}{X - t(x)},$$

and engages with the verifier in an IOP for showing that

$$\sum_{x \in H} \sum_{i=0}^{M-1} \frac{v_T(X)}{X + f_i(x)} = R_T(X). \quad (47)$$

The verifier queries $v_T(X)$ and $R_T(X)$ at a random challenge $\alpha \leftarrow_{\$} F$, and both prover and verifier run a sumcheck argument for

$$\sum_{x \in H} \sum_{i=0}^{M-1} \frac{1}{\varphi_i(x)} = \frac{R_T(\alpha)}{v_T(\alpha)},$$

where $\varphi_i(x) = \alpha + f_i(x)$. For this the prover provides the Lagrange representation of the sumcheck polynomial $\phi(X)$ subject to the domain identity

$$\phi(g \cdot x) - \phi(x) + \frac{R_T(\alpha)}{|H| \cdot v_T(\alpha)} = \sum_{i=0}^{M-1} \frac{1}{\varphi_i(x)}$$

for all $x \in H$, or as the polynomial identity

$$\left(\phi(g \cdot X) - \phi(X) + \frac{R_T(\alpha)}{|H| \cdot v_T(\alpha)} \right) \cdot \prod_{i=0}^{M-1} \varphi_i(X) = \prod_{i=0}^{M-1} \varphi_i(X) \cdot \sum_{i=0}^{M-1} \frac{1}{\varphi_i(X)} \mod v_H(X). \quad (48)$$

The overall identity is of the form

$$Q(\phi_0(X), \dots, \phi_{M-1}(X), \phi(X), \phi(g \cdot X)) = 0 \mod v_H(X),$$

with Q having $\nu = M + 2$ variables and absolute degree $d = M + 1$. The prover provides the Lagrange representation of the component polynomials $q_1(X), \dots, q_M(X)$ for the overall quotient. The verifier samples $\beta \leftarrow_{\$} F$ and checks (48) at $X = \beta$ by quering the involved polynomials at the needed points.

Cost analysis

Let us count the number of field operations of the oracle prover. Computing the coefficients of $R_T(X)$ by interpolating the multiplicities m as a function on the table range $T = \{t(x) : x \in H\}$ costs

$$\frac{|H|}{8} \left(3 \cdot \log^2 |H| + 11 \cdot \log |H| + 3 \right)$$

field multiplications, using pre-computed intermediate products of the table factors $X + t(x)$ (both the coefficients and their values used for obtaining them). See [GK22] and the references therein for details on

⁸In a variant of the protocol, communicated by A. Gabizon, the prover provides the Lagrange representation of $R_T(X)$ with respect to the Lagrange basis of the table set $\{t(x) : x \in H\}$. As Lagrange IOPs with respect to several Lagrange bases have different impacts on the Lagrange commitment scheme (e.g., it leads to a separate commitment in a KZG-like scheme, or a separate inner product argument for an IPA-like scheme), we do not compare with this variant.

Lagrange interpolation over general sets. The Lagrange representation of $R_T(X)$ over H is then obtained by another

$$\frac{|H|}{2} \cdot \log |H|$$

multiplications. The values of the reciprocals $h(x) = \frac{1}{\alpha + f(x)}$ over H cost $3 \cdot |H|$ multiplications (using batch inversion), the computation of ϕ does not involve a single field multiplication. Using the fractional representation (48) for evaluating Q using batch inversion yields the equivalent multiplicative complexity

$$Q_M = M + 1 + 3 \cdot M = 4 \cdot M + 1,$$

and by (34) computing the components q_1, \dots, q_M of the overall quotient costs

$$|H| \cdot (M + 1) \left(\frac{1}{2} \cdot \log |H| + \frac{(M + 1)}{2} \cdot \log((M + 1) \cdot |H|) \right) + (M + 1) \cdot |H| \cdot (4 \cdot M + 1).$$

The overall cost of the oracle prover is

$$\begin{aligned} |H| \cdot \left((M + 1) \left(\frac{M + 2}{2} \cdot \log |H| + \frac{(M + 1)}{2} \cdot \log(M + 1) \right) + (M + 1) \cdot (4 \cdot M + 1) \right. \\ \left. + \frac{3}{8} \cdot \log^2 |H| + \frac{15}{8} \cdot \log |H| + \frac{3}{8} + 5 \right), \end{aligned} \quad (49)$$

whereas the oracle costs are equivalent to $M + 2$ domain evaluations of size $|H|$, corresponding to R_T , ϕ , and q_1, \dots, q_M .

A.4 Comparison

Based on the equivalent number of field multiplications from Table 1, we obtain the following estimation of the logarithmic derivative lookups over plookup, as the ratio of their number of field multiplications $r = \mathbb{M}(\text{Plookup})/\mathbb{M}(\log D)$: For all domain sizes $\log |H| = 12, 14, 16, 18$ we get a maximum advantage of about $r_{\max} \approx 1.8$, whereas

$$r = \mathbb{M}(\text{plookup})/\mathbb{M}(\log D) \geq 1.75$$

whenever $M \geq 12$. Although having worse asymptotic complexity, the flookup strategy outperforms the logarithmic derivative plookup at these domain sizes, see Table ???. Its maximum advantage is about 1.25 – 1.3 is throughout at $M = 1$, whereas that advantage tends to 1 at high number of columns (for $M \geq 33$ it is less than 5%).